

Attention in NLP: Visualizing the Words that Affect Sentiment Predictions

GitHub Link: <https://github.com/Kurakulagopi23/Attension-in-NLP>

Student Name : Gopi Kurakula

Student ID : 24142071

1. Introduction

1.1 Sentiment Classification in Natural Language Processing

Sentiment classification is a prevalent job in Natural Language Processing wherein a model determines if a text conveys a positive or negative sentiment. It is utilized in domains such as film critique assessment, product evaluation, customer service oversight, and social media analysis. This tutorial demonstrates sentiment categorization via IMDB movie reviews, with each review categorized as either positive or negative.

In contrast to numerical or visual input, text cannot be utilized directly by a neural network. A sentence has words, which must be transformed into numerical representations prior to processing. The sequence of words is significant. The phrases “not good” and “good” possess distinct meanings, while both using the term “good.” NLP differs from basic tabular categorization as the model must comprehend both the semantics of words and their sequential arrangement.

1.2 Tutorial Focus

This tutorial centers on an attention-based neural network for sentiment categorization. The model employs four critical stages: tokenization, padding, word embeddings, and sequence learning utilizing a Bidirectional Long Short-Term Memory network. An attention layer is subsequently incorporated to allocate significance weights to various word places within the review.

The primary objective of employing attention is not solely to classify sentiment, but also to enhance the model's inspectability. A conventional sequence model can generate a forecast without indicating the words that impacted it. An attention mechanism offers a straightforward method to visualize the significance of words. This renders the technique advantageous for instruction and fundamental model interpretation.

2. Text Preparation for a Neural Network

2.1 Tokenization

A neural network is incapable of directly comprehending unprocessed text.

Tokenization transforms words into numerical identifiers. The line "this movie was surprisingly good" can be expressed as a numerical sequence. Every numeral is associated with a term in the lexicon. The number itself lacks intrinsic value; it merely serves as a reference enabling the model to identify a learned representation for that word.

The IMDB dataset contains reviews encoded as sequences of integers. This renders the dataset appropriate for a brief tutorial, as the emphasis can remain on comprehending the NLP model rather than constructing a comprehensive text cleaning pipeline.

2.2 Padding

Movie reviews vary in length. Some assessments are concise, but others are considerably more extensive. Neural networks often handle batches of uniformly sized inputs, necessitating that reviews be padded or shortened to a predetermined length. Each review in the notebook is restricted to 80 tokens. Concise reviews are supplemented with zeros, whilst extensive reviews are truncated to the maximum length.

Figure 1: An illustration of tokenization and padding, demonstrating the conversion of a brief review into word indices and its adjustment to a predetermined sequence length.

Figure 1: Tokenisation and Padding Example

```
Original text:
"this movie was surprisingly good"

Tokenised sequence:
[45, 17, 23, 981, 49]

Padded sequence:
[45, 17, 23, 981, 49, 0, 0, 0]
```

Alt text: A textual diagram illustrating an original movie review sentence, its tokenized integer sequence, and the zero-padded variant.

2.3 Overview of Preprocessing Procedures

Step	Purpose
Tokenisation	Converts words into integer indexes.
Padding	Makes all review sequences the same length.

Truncation	Shortens very long reviews to the selected maximum length.
Embedding lookup	Converts word indexes into dense vectors.

The preprocessing stages are crucial as they convert raw text into a format suitable for neural network processing.

3. Embeddings and Sequential Learning

3.1 Word Embeddings

Subsequent to tokenization, each word is just represented by an integer index. This is insufficient for comprehending meaning. An embedding layer transforms each word index into a dense numerical vector. These vectors are acquired throughout the training process. Lexical items that occur in analogous settings may cultivate comparable representations, enabling the model to discern valuable sentiment patterns.

For instance, terms like “excellent,” “brilliant,” and “enjoyable” may serve as effective indicators of positive sentiment, whereas terms such as “boring,” “awful,” and “disappointing” may function as reliable indicators of negative sentiment. The embedding layer does not explicitly retain these significances. Rather, it acquires word representations through the review categorization job.

3.2 The Significance of Word Order

The meaning of text is heavily influenced by word arrangement. The bag-of-words methodology can identify individual words; nevertheless, it fails to accurately represent their sequence. The expression “the film was not bad” differs from “the film was bad.” The placement of “not” alters the phrase's meaning. Consequently, a sequential model is required to analyze words in a specific order.

Long Short-Term Memory networks are engineered for sequential data processing. They sequentially process information and retain data from prior states. This renders them advantageous for writing, because the significance of a word frequently relies on adjacent words.

3.3 Bidirectional Long Short-Term Memory

A Bidirectional LSTM enhances sequence learning by processing text in both directions. The forward LSTM processes the review from start to finish, whereas the reverse LSTM processes it from finish to start. The results from both directions are amalgamated. This provides the model with access to both preceding and

subsequent information when encoding each word location.

In mood categorization, this is advantageous as the interpretation of a sentence may rely on the preceding and succeeding words. A BiLSTM can consequently provide more nuanced word-level representations than a unidirectional sequence model.

4. Attention Mechanism in Natural Language Processing

4.1 Definition of Attention

Attention enables the model to allocate varying significance levels to distinct word positions. In sentiment classification, not all words contribute equally to the ultimate prediction. Terms like “excellent,” “terrible,” “boring,” or “brilliant” frequently convey greater passion than commonplace terms such as “the,” “and,” or “movie.” Attention enables the model to construct a sentence representation by prioritizing more informative positions.

4.2 Mathematical demonstration

Upon processing a review, the BiLSTM generates a hidden representation for each word position:

$$h_1, h_2, h_3, \dots, h_T$$

In this context, (T) denotes the sequence length, whereas (h_t) signifies the BiLSTM output at the word position (t). The attention layer computes a score for every word position:

$$e_t = \tanh(W_a h_t + b_a)$$

where (W_a) and (b_a) represent adjustable attention parameters. The ratings are subsequently transformed into attention weights via softmax.

$$\alpha_t = \frac{\exp(e_t)}{\sum_{j=1}^T \exp(e_j)}$$

The attention weights total to 1, functioning as a distribution of significance among the words in the review. An elevated (α_t) indicates that the model assigns greater significance to the word position (t). The ultimate context vector is computed as a weighted summation:

$$c = \sum_{t=1}^T \alpha_t h_t$$

The context vector (c) serves as the review representation utilized by the classifier. The ultimate emotion prediction relies on a weighted amalgamation of word-level information rather than treating all words identically.

4.3 Sentiment Forecasting


Subsequent to the generation of the context vector by the attention layer, the classifier transforms it into a probability.

$$\hat{y} = \sigma(Wc + b)$$

where (σ) denotes the sigmoid function. A score approaching 1 indicates good attitude, whereas a value nearing 0 signifies negative sentiment. The attention weights enhance interpretability by indicating which word positions significantly influenced the context vector.

Figure 2: Attention-driven NLP pipeline from review text to sentiment analysis.

Attention-based NLP Pipeline



A flow diagram illustrating the progression of review text through tokenization, padding, embedding, BiLSTM, attention mechanisms, and culminating in sentiment prediction.

5. Constructing the BiLSTM-Attention Classifier

5.1 Architectural Framework

The attention-based sentiment classifier adheres to a certain arrangement of layers. The input consists of a padded review depicted as a series of word indices. The embedding layer transforms each index into a dense vector. The Bidirectional LSTM processes the embedded sequence from both directions, generating a contextual representation for each word position. The attention layer receives these word-level representations and generates a weighted sentence representation. A dense classifier ultimately transforms this representation into a sentiment prediction.

Table 1: Layer structure of the BiLSTM-attention sentiment classifier, showing the function of each layer from padded text input to final sentiment prediction.

Layer	Function
Input layer	Receives the padded review sequence.
Embedding layer	Converts word indexes into dense word vectors.
Bidirectional LSTM	Learns contextual sequence representations.
Attention layer	Assigns importance weights to word positions.
Dense layer	Learns higher-level sentiment patterns.
Sigmoid output	Produces a positive or negative sentiment probability.

This architecture is efficient for a tutorial as each phase has a distinct function. The embedding layer manages word representation, the BiLSTM addresses word order and context, and attention facilitates interpretation by emphasizing significant word-level contributions.

5.2 Binary Sentiment Output

The terminal output layer employs a sigmoid activation function due to the binary nature of the task, which encompasses two potential labels: positive and negative. The result is a probability value ranging from 0 to 1. A score approaching 1 indicates good attitude, whereas a value nearing 0 signifies negative sentiment. The model is trained utilizing binary cross-entropy loss:

$$L = -[y\log(\hat{y}) + (1 - y)\log(1 - \hat{y})]$$

In this context, y represents the actual sentiment label while $\hat{y}$ denotes the predicted sentiment label.


This represents the anticipated probability. This loss function imposes a greater penalty on confident

6. Exhibition and Visual Analysis

6.1 Training Conduct

The notebook employs a limited portion of the IMDB dataset and trains the model for two epochs. This maintains the code's efficiency for a tutorial while effectively illustrating the model's behavior. During the demonstration, training accuracy rose from roughly 0.54 to 0.74, but validation accuracy ascended from about 0.59 to 0.74. These values indicate that the machine has acquired valuable sentiment patterns from the review data.

Figure 3: Training and validation accuracy of the BiLSTM-attention sentiment classifier.




This result does not assert that the model is completely optimized. An expanded training dataset, increased epochs, and meticulous tuning may enhance performance. The outcome mostly serves to reinforce the instructional explanation: the model is capable of learning a sentiment classification task while simultaneously generating attention weights for examination.

6.2 Visualization of Attention

The attention chart illustrates the terms that garnered more attention weights in a specific review. Attention weights, generated by a softmax function, assign a value to each word position that contributes to the final context vector. Longer bars denote word positions that exerted a more significant impact on the sentence representation utilized for prediction.

Figure 4: Attention weights for certain terms in a review.



The attention visualization must be interpreted with caution. In the notebook output, certain high-attention terms are plainly pertinent, but others may be less apparent. This is anticipated in a brief demonstration utilizing a compact model. Attention serves as a valuable diagnostic instrument; yet, it does not offer a comprehensive elucidation of the model's rationale. It is more accurately perceived as an acquired weighing mechanism that elucidates aspects of the model's decision-making process.

6.3 Examples of Predictions

The prediction examples link the model's results to comprehensible text. Each sample comprises a review excerpt, the actual label, the anticipated label, and identified high-attention terms.

Table 2: Prediction examples from the BiLSTM-attention model

Review type	True label	Predicted label	Example high-attention words
Negative review	Negative	Negative	miss, wasn't, give
Positive review	Positive	Positive	movie, being, disgusted
Positive review	Positive	Positive	train, concept, getting

These examples illustrate both the merits and constraints of attentiveness. The model can generate accurate predictions; nevertheless, the chosen terms must be subjected to critical evaluation. Certain words may garner significant attention due to their placement or context rather than their explicit emotional connotation.

7. Practical Recommendations and Constraints

7.1 The Utility of Attention

Attention is beneficial when a sequence model must deliver more than a singular prediction. In sentiment classification, it aids in determining which segments of a review impacted the sentence representation. This is beneficial for instruction, troubleshooting, and fundamental analysis. It can also facilitate error analysis by indicating whether the model is concentrating on sentiment-related terms or extraneous words.

7.2 Prevailing Limitations

The primary restriction is that attention does not provide a comprehensive explanation. A high attention score indicates that a word's position significantly influenced the context vector; nevertheless, it does not demonstrate that the model reasoned similarly to a human reader. The subsequent limitation is padding. If padding is not managed meticulously, the model may derive insights from bogus zero tokens. The third constraint pertains to sequence length. Limiting reviews to 80 tokens accelerates the tutorial, however it may eliminate valuable information from lengthier evaluations.

Another constraint is dataset bias. IMDB evaluations exhibit linguistic patterns characteristic of film critiques. A model trained on this data should not be presumed to perform uniformly across product reviews, political remarks, medical documentation, or job evaluations.

7.3 Potential Enhancements

The lesson model can be enhanced by utilizing additional training samples, extending the training duration, and implementing more robust masking for padded tokens. An evaluation against a standard LSTM model could elucidate the impact of attention more distinctly. While more sophisticated tutorials may juxtapose this model with transformer-based attention, the BiLSTM-attention framework serves as a more accessible introductory example of attention-based NLP.

8. Ethical Considerations

Sentiment algorithms may err when analyzing text that includes sarcasm, humor, slang, dialect, or mixed sentiments. A review stating “great, another boring film” may pose challenges due to the sarcastic usage of the positive term “great.” Text models may also exhibit bias derived from their training data. If specific writing styles, dialects, or cultural expressions are inadequately represented, the model may interpret them with diminished accuracy.

In low-risk movie review classification, these errors are predominantly instructional. In high-risk environments, such as staff surveillance, customer profiling, or moderation systems, erroneous sentiment predictions may yield unjust repercussions. Attention visualization can facilitate examination; however, it should not serve as the sole rationale for automated judgments.

9. Conclusion

Attention enhances the interpretability of neural NLP models by allocating weights to word positions. This tutorial involved the tokenization, padding, embedding, processing of IMDB reviews using a BiLSTM, and subsequent passage through an attention layer. The attention technique generated a weighted representation of the sentence prior to sentiment prediction. The demonstration illustrated the visualization and interpretation of attention weights, while also emphasizing their limitations. Attention is advantageous since it links sequence learning with word-level examination, hence enhancing the comprehensibility and analysis of sentiment classifiers.

References

Bahdanau, D., Cho, K. and Bengio, Y. (2015) 'Neural machine translation by jointly learning to align and translate', *International Conference on Learning Representations*. Available at: arXiv:1409.0473.

Chollet, F. (2021) *Deep Learning with Python*. 2nd edn. Shelter Island: Manning.

Hochreiter, S. and Schmidhuber, J. (1997) 'Long short-term memory', *Neural Computation*, 9(8), pp. 1735–1780. doi: 10.1162/neco.1997.9.8.1735.

Keras (2024) 'IMDB movie review sentiment classification dataset'. Available at: Keras API documentation.

Maas, A.L., Daly, R.E., Pham, P.T., Huang, D., Ng, A.Y. and Potts, C. (2011) 'Learning word vectors for sentiment analysis', *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, pp. 142–150.

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, L. and Polosukhin, I. (2017) 'Attention is all you need', *Advances in Neural Information Processing Systems*, 30.